

```
!pip install numpy
!pip install matplotlib
```

```
Requirement already satisfied: numpy in
/usr/local/lib/python3.10/dist-packages (1.23.5)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.1.1)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (4.43.1)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)
Requirement already satisfied: numpy>=1.20 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (1.23.5)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (23.2)
Requirement already satisfied: pillow>=6.2.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.1)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.10/dist-packages (from matplotlib) (2.8.2)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7-
>matplotlib) (1.16.0)
```

```
import pandas as pd
import numpy as np
import os
```

```
# Get the current directory
current_directory = os.getcwd()
```

```
# Define the CSV file's name
csv_file_name = "Regression.csv"
```

```
# Construct the full path to the CSV file
csv_file_path = os.path.join(current_directory, csv_file_name)
```

```
# Load the dataset from CSV file
data = pd.read_csv(csv_file_name)
```

```
data.columns
data
```

	station	Date	Present_Tmax	Present_Tmin	LDAPS_RHmin	\
0	1.0	2013-06-30	28.7	21.4	58.255688	

1	2.0	2013-06-30	31.9	21.6	52.263397
2	3.0	2013-06-30	31.6	23.3	48.690479
3	4.0	2013-06-30	32.0	23.4	58.239788
4	5.0	2013-06-30	31.4	21.9	56.174095
...	...	...	...	...	...
7747	23.0	2017-08-30	23.3	17.1	26.741310
7748	24.0	2017-08-30	23.3	17.7	24.040634
7749	25.0	2017-08-30	23.2	17.4	22.933014
7750	NaN	NaN	20.0	11.3	19.794666
7751	NaN	NaN	37.6	29.9	98.524734

	LDAPS_RHmax	LDAPS_Tmax_lapse	LDAPS_Tmin_lapse	LDAPS_WS
LDAPS_LH \				
0	91.116364	28.074101	23.006936	6.818887
69.451805				
1	90.604721	29.850689	24.035009	5.691890
51.937448				
2	83.973587	30.091292	24.565633	6.138224
20.573050				
3	96.483688	29.704629	23.326177	5.650050
65.727144				
4	90.155128	29.113934	23.486480	5.735004
107.965535				

...	...	...	...	...
...				
7747	78.869858	26.352081	18.775678	6.148918
72.058294				
7748	77.294975	27.010193	18.733519	6.542819
47.241457				
7749	77.243744	27.939516	18.522965	7.289264
9.090034				
7750	58.936283	17.624954	14.272646	2.882580
13.603212				
7751	100.000153	38.542255	29.619342	21.857621
213.414006				

	...	LDAPS_PPT2	LDAPS_PPT3	LDAPS_PPT4	lat	lon
DEM \						
0	...	0.000000	0.000000	0.000000	37.6046	126.991
212.3350						
1	...	0.000000	0.000000	0.000000	37.6046	127.032
44.7624						
2	...	0.000000	0.000000	0.000000	37.5776	127.058
33.3068						
3	...	0.000000	0.000000	0.000000	37.6450	127.022
45.7160						
4	...	0.000000	0.000000	0.000000	37.5507	127.135
35.0380						
...	...	...	...	...	...	...

```

...
7747 ... 0.000000 0.000000 0.000000 37.5372 126.891
15.5876
7748 ... 0.000000 0.000000 0.000000 37.5237 126.909
17.2956
7749 ... 0.000000 0.000000 0.000000 37.5237 126.970
19.5844
7750 ... 0.000000 0.000000 0.000000 37.4562 126.826
12.3700
7751 ... 21.621661 15.841235 16.655469 37.6450 127.135
212.3350

```

	Slope	Solar radiation	Next_Tmax	Next_Tmin
0	2.785000	5992.895996	29.1	21.2
1	0.514100	5869.312500	30.5	22.5
2	0.266100	5863.555664	31.1	23.9
3	2.534800	5856.964844	31.7	24.3
4	0.505500	5859.552246	31.2	22.5
...	...	...	...	...
7747	0.155400	4443.313965	28.3	18.1
7748	0.222300	4438.373535	28.6	18.8
7749	0.271300	4451.345215	27.8	17.4
7750	0.098475	4329.520508	17.4	11.3
7751	5.178230	5992.895996	38.9	29.8

```
[7752 rows x 25 columns]
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7752 entries, 0 to 7751
```

```
Data columns (total 25 columns):
```

#	Column	Non-Null	Count	Dtype
0	station	7750	non-null	float64
1	Date	7750	non-null	object
2	Present_Tmax	7682	non-null	float64
3	Present_Tmin	7682	non-null	float64
4	LDAPS_RHmin	7677	non-null	float64
5	LDAPS_RHmax	7677	non-null	float64
6	LDAPS_Tmax_lapse	7677	non-null	float64
7	LDAPS_Tmin_lapse	7677	non-null	float64
8	LDAPS_WS	7677	non-null	float64
9	LDAPS_LH	7677	non-null	float64
10	LDAPS_CC1	7677	non-null	float64
11	LDAPS_CC2	7677	non-null	float64
12	LDAPS_CC3	7677	non-null	float64
13	LDAPS_CC4	7677	non-null	float64
14	LDAPS_PPT1	7677	non-null	float64
15	LDAPS_PPT2	7677	non-null	float64

16	LDAPS_PPT3	7677	non-null	float64
17	LDAPS_PPT4	7677	non-null	float64
18	lat	7752	non-null	float64
19	lon	7752	non-null	float64
20	DEM	7752	non-null	float64
21	Slope	7752	non-null	float64
22	Solar radiation	7752	non-null	float64
23	Next_Tmax	7725	non-null	float64
24	Next_Tmin	7725	non-null	float64

dtypes: float64(24), object(1)

memory usage: 1.5+ MB

*# Remove the first two columns (station and date) and select columns 3-23 as features*

data = data.iloc[:, 2:23]

data.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 7752 entries, 0 to 7751

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
0	Present_Tmax	7682 non-null	float64
1	Present_Tmin	7682 non-null	float64
2	LDAPS_RHmin	7677 non-null	float64
3	LDAPS_RHmax	7677 non-null	float64
4	LDAPS_Tmax_lapse	7677 non-null	float64
5	LDAPS_Tmin_lapse	7677 non-null	float64
6	LDAPS_WS	7677 non-null	float64
7	LDAPS_LH	7677 non-null	float64
8	LDAPS_CC1	7677 non-null	float64
9	LDAPS_CC2	7677 non-null	float64
10	LDAPS_CC3	7677 non-null	float64
11	LDAPS_CC4	7677 non-null	float64
12	LDAPS_PPT1	7677 non-null	float64
13	LDAPS_PPT2	7677 non-null	float64
14	LDAPS_PPT3	7677 non-null	float64
15	LDAPS_PPT4	7677 non-null	float64
16	lat	7752 non-null	float64
17	lon	7752 non-null	float64
18	DEM	7752 non-null	float64
19	Slope	7752 non-null	float64
20	Solar radiation	7752 non-null	float64

dtypes: float64(21)

memory usage: 1.2 MB

*# Drop missing values (NaNs)*

data.dropna(inplace=True)

def min\_max\_scaling(data):

```

min_val = data.min()
max_val = data.max()
scaled_data = (data - min_val) / (max_val - min_val)
return scaled_data

# Min-Max scaling
for col in data.columns:
    data[col] = min_max_scaling(data[col])

data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 7607 entries, 0 to 7751
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Present_Tmax          7607 non-null   float64
1   Present_Tmin          7607 non-null   float64
2   LDAPS_RHmin           7607 non-null   float64
3   LDAPS_RHmax           7607 non-null   float64
4   LDAPS_Tmax_lapse      7607 non-null   float64
5   LDAPS_Tmin_lapse      7607 non-null   float64
6   LDAPS_WS              7607 non-null   float64
7   LDAPS_LH              7607 non-null   float64
8   LDAPS_CC1             7607 non-null   float64
9   LDAPS_CC2             7607 non-null   float64
10  LDAPS_CC3             7607 non-null   float64
11  LDAPS_CC4             7607 non-null   float64
12  LDAPS_PPT1            7607 non-null   float64
13  LDAPS_PPT2            7607 non-null   float64
14  LDAPS_PPT3            7607 non-null   float64
15  LDAPS_PPT4            7607 non-null   float64
16  lat                   7607 non-null   float64
17  lon                   7607 non-null   float64
18  DEM                   7607 non-null   float64
19  Slope                 7607 non-null   float64
20  Solar radiation       7607 non-null   float64
dtypes: float64(21)
memory usage: 1.3 MB

def gradient_descent(X, y, theta, alpha, num_iterations):
    m = len(y)
    for iteration in range(num_iterations):
        gradients = -2/m * X.T @ (y - X @ theta)
        theta -= alpha * gradients
    return theta

# Number of trials
num_trials = 10

```

```

# Initialize arrays to store RMSE values
mean_training_rmse_tmax = []
mean_testing_rmse_tmax = []
mean_training_rmse_tmin = []
mean_testing_rmse_tmin = []

import matplotlib.pyplot as plt

# Repeat the process for 10 trials
for trial in range(num_trials):
    print(f"Trial {trial + 1}:")

    # Shuffle the data to ensure randomness
    data = data.sample(frac=1).reset_index(drop=True)

    # Split the data into training and testing sets (80% training, 20%
testing)
    train_size = int(0.8 * len(data))
    training_data = data[:train_size]
    testing_data = data[train_size:]

    # Extract features and target variables
    X_train = training_data.iloc[:, :-2].values
    y_train = training_data.iloc[:, -2:].values
    X_test = testing_data.iloc[:, :-2].values
    y_test = testing_data.iloc[:, -2:].values

    # Initialize theta with zeros
    theta = np.zeros((X_train.shape[1], 2))

    # Define hyperparameters
    alpha = 0.01 # Learning rate
    num_iterations = 1000 # Number of iterations

    # Initialize arrays to store error values during training
    training_errors_tmax = []
    testing_errors_tmax = []
    training_errors_tmin = []
    testing_errors_tmin = []

    for i in range(num_iterations):
        # Perform gradient descent to train the linear regression
model
        theta = gradient_descent(X_train, y_train, theta, alpha, 1)

        # Make predictions using the trained model
        predicted_train = X_train @ theta
        predicted_test = X_test @ theta

        # Calculate RMSE for TMax and TMin separately

```

```

        training_rmse_tmax = np.sqrt(np.mean((y_train[:, 0] -
predicted_train[:, 0]) ** 2))
        testing_rmse_tmax = np.sqrt(np.mean((y_test[:, 0] -
predicted_test[:, 0]) ** 2))

        training_rmse_tmin = np.sqrt(np.mean((y_train[:, 1] -
predicted_train[:, 1]) ** 2))
        testing_rmse_tmin = np.sqrt(np.mean((y_test[:, 1] -
predicted_test[:, 1]) ** 2))

        # Append RMSE values to the arrays
        training_errors_tmax.append(training_rmse_tmax)
        testing_errors_tmax.append(testing_rmse_tmax)
        training_errors_tmin.append(training_rmse_tmin)
        testing_errors_tmin.append(testing_rmse_tmin)

    print(f"Training RMSE for TMax: {training_rmse_tmax}, Testing RMSE
for TMax: {testing_rmse_tmax}")
    print(f"Training RMSE for TMin: {training_rmse_tmin}, Testing RMSE
for TMin: {testing_rmse_tmin}")

    # Append RMSE values to the lists for each trial
    mean_training_rmse_tmax.append(training_rmse_tmax)
    mean_testing_rmse_tmax.append(testing_rmse_tmax)
    mean_training_rmse_tmin.append(training_rmse_tmin)
    mean_testing_rmse_tmin.append(testing_rmse_tmin)

    # Plot error curves for this trial
    plt.figure()
    plt.subplot(2, 1, 1)
    plt.plot(range(num_iterations), training_errors_tmax,
label="Training TMax")
    plt.plot(range(num_iterations), testing_errors_tmax,
label="Testing TMax")
    plt.xlabel("Iterations")
    plt.ylabel("RMSE")
    plt.legend()
    plt.title("TMax Error Curves")
    # Show all the error curve plots
    plt.show()

    plt.subplot(2, 1, 2)
    plt.plot(range(num_iterations), training_errors_tmin,
label="Training TMin")
    plt.plot(range(num_iterations), testing_errors_tmin,
label="Testing TMin")
    plt.xlabel("Iterations")
    plt.ylabel("RMSE")
    plt.legend()
    plt.title("TMin Error Curves")

```

```
# Show all the error curve plots
plt.show()
```



```
# Calculate and print the mean RMSE over all trials
```

```
mean_training_rmse_tmax = np.mean(mean_training_rmse_tmax)
mean_testing_rmse_tmax = np.mean(mean_testing_rmse_tmax)
mean_training_rmse_tmin = np.mean(mean_training_rmse_tmin)
mean_testing_rmse_tmin = np.mean(mean_testing_rmse_tmin)
print(f"Mean Training RMSE for TMax: {mean_training_rmse_tmax}, Mean  
Testing RMSE for TMax: {mean_testing_rmse_tmax}")
print(f"Mean Training RMSE for TMin: {mean_training_rmse_tmin}, Mean  
Testing RMSE for TMin: {mean_testing_rmse_tmin}")
```

Trial 1:

Training RMSE for TMax: 0.17063527646260018, Testing RMSE for TMax:  
0.1694633642162574

Training RMSE for TMin: 0.24723635215505882, Testing RMSE for TMin:  
0.24196286868984512


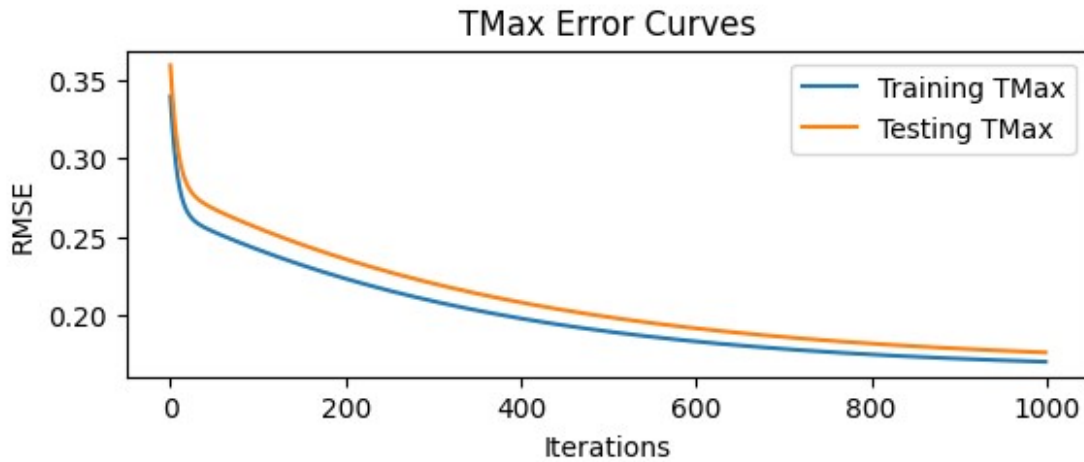




Trial 2:

Training RMSE for TMax: 0.17004381724547424, Testing RMSE for TMax: 0.17598605109466572



Training RMSE for TMin: 0.24658145878505758, Testing RMSE for TMin: 0.24458585770102015




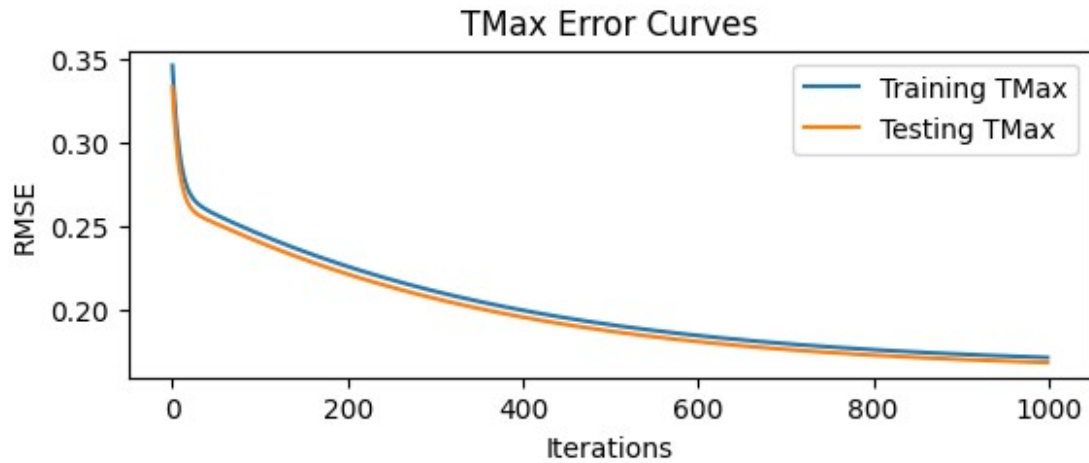
Trial 3:

Training RMSE for TMax: 0.17000677126661776, Testing RMSE for TMax: 0.1748772491004664


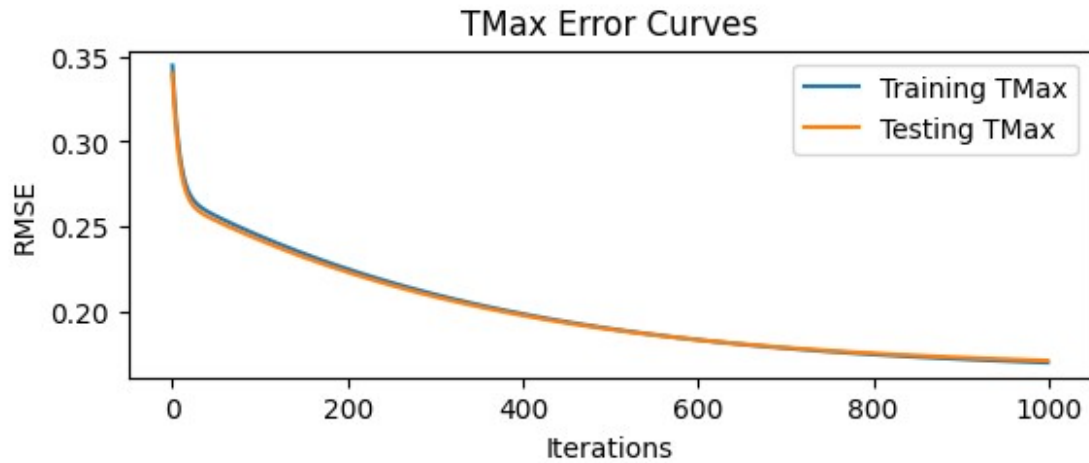
Training RMSE for TMin: 0.24686987562473395, Testing RMSE for TMin: 0.24373692871286998



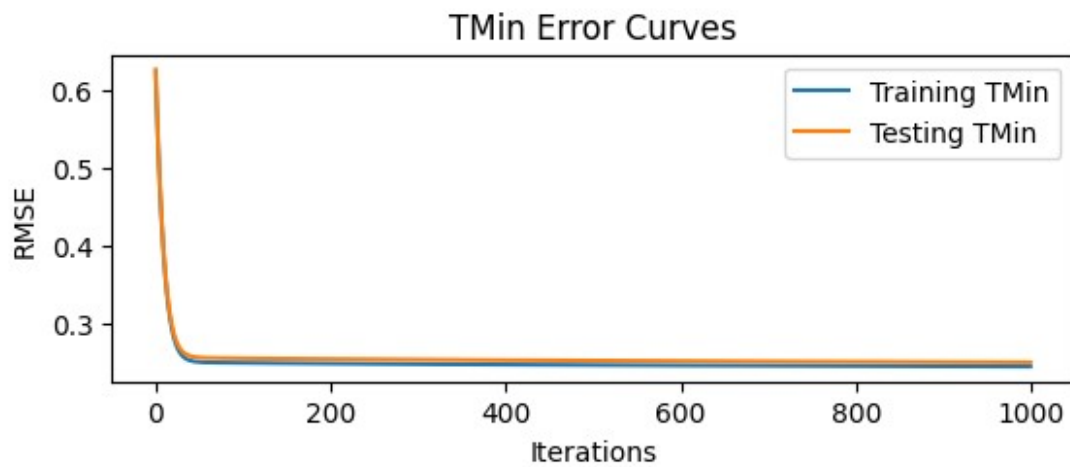

Trial 4:  
Training RMSE for TMax: 0.17098692030689933, Testing RMSE for TMax: 0.1680057400455115  
Training RMSE for TMin: 0.24678694130406179, Testing RMSE for TMin: 0.24381092371211863





Trial 5:  
Training RMSE for TMax: 0.17012847646832308, Testing RMSE for TMax: 0.17127991508309567  
Training RMSE for TMin: 0.2460622820538406, Testing RMSE for TMin: 0.24730660096866067





Trial 6:  
Training RMSE for TMax: 0.17101418765328733, Testing RMSE for TMax: 0.17038899723213533  
Training RMSE for TMin: 0.24527028047275387, Testing RMSE for TMin: 0.2506363394671873





Trial 7:  
Training RMSE for TMax: 0.17084766529112988, Testing RMSE for TMax: 0.1671296999832878  
Training RMSE for TMin: 0.24503321419379276, Testing RMSE for TMin: 0.2513254925242216



Trial 8:  
Training RMSE for TMax: 0.17097758218902176, Testing RMSE for TMax: 0.16873270277782987  
Training RMSE for TMin: 0.24670919813658798, Testing RMSE for TMin: 0.24399437401744853

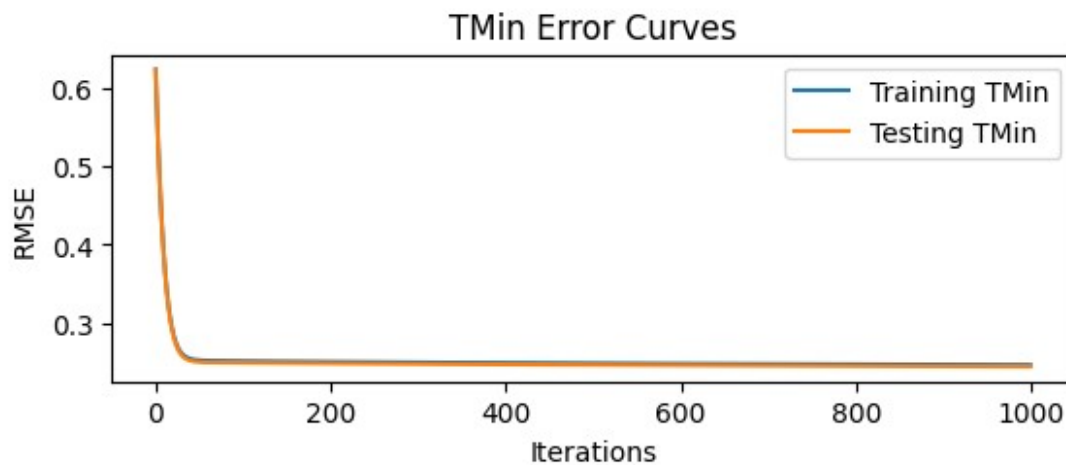



Trial 9:  
Training RMSE for TMax: 0.170155933057318, Testing RMSE for TMax: 0.17063426118530853  
Training RMSE for TMin: 0.24704636419614837, Testing RMSE for TMin: 0.24256783042204463



Trial 10:  
Training RMSE for TMax: 0.17051287117039787, Testing RMSE for TMax: 0.17254713396319155  
Training RMSE for TMin: 0.24668986863814624, Testing RMSE for TMin: 0.2445582843801485





Mean Training RMSE for TMax: 0.17053095011110694, Mean Testing RMSE for TMax: 0.17090451146817495  
Mean Training RMSE for TMin: 0.24642858355601818, Mean Testing RMSE for TMin: 0.24544855005955651

- **Impact of Learning Rate:** The learning rate is a crucial hyperparameter in gradient descent. Setting it too high can result in overshooting the minimum of the cost function and lead to divergence. Setting it too low can result in slow convergence. It's important to experiment with different learning rates to find the one that results in the fastest convergence without divergence.
- **Impact of Data Preprocessing:** Data preprocessing, including handling missing values and normalization, can have a significant impact on prediction performance. For example, removing rows with missing values may lead to loss of valuable data, while various normalization techniques can affect the scale and distribution of features. It's essential to carefully choose preprocessing methods based on the characteristics of your data.

- **Impact of Random Initialization:** The initial random weights can significantly affect the training process. In the case of linear regression, the initial values of  $\theta$  can affect convergence speed. Random initialization of  $\theta$  can lead to different results in different runs. Experimenting with different random seeds can provide insights into the stability and reproducibility of the model.
- **Impact of the Number of Iterations:** The number of iterations in gradient descent is another hyperparameter. Too few iterations may result in an underfitted model, while too many iterations may lead to overfitting. Monitoring the cost function's convergence and early stopping when it stops improving can save computation time.